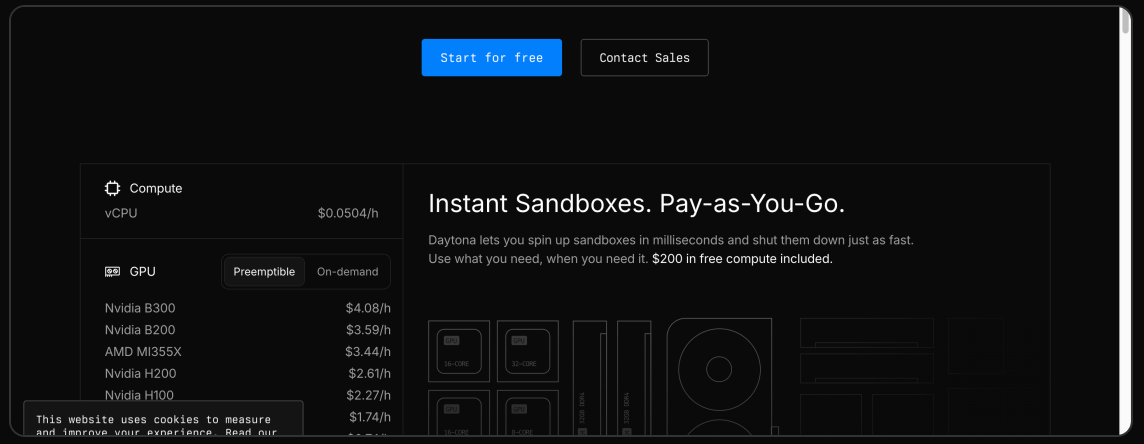
AgentReadable

# 10/10 wrong price

"I need an RTX 4090 that won't be interrupted. How much per hour?"
gpt-oss 20B, reading a reproduction of this page in a real browser, answered **$0.57** every time. The on-demand price is **$0.99**.

`github.com/TaeWoo1/daytona-nosana-hackathon` · Daytona · Nosana




daytona.io/pricing, captured 2026-10-18 (data/raw/ref). GPU table opens on **Preemptible**.

gpt-oss 20B, layout A, q1, 10 runs: confident-wrong 10/10, clicked the tab 0/10. Source: results/summary.json → behavior (runs_rendered_v2_n10).

# Human view vs AI view



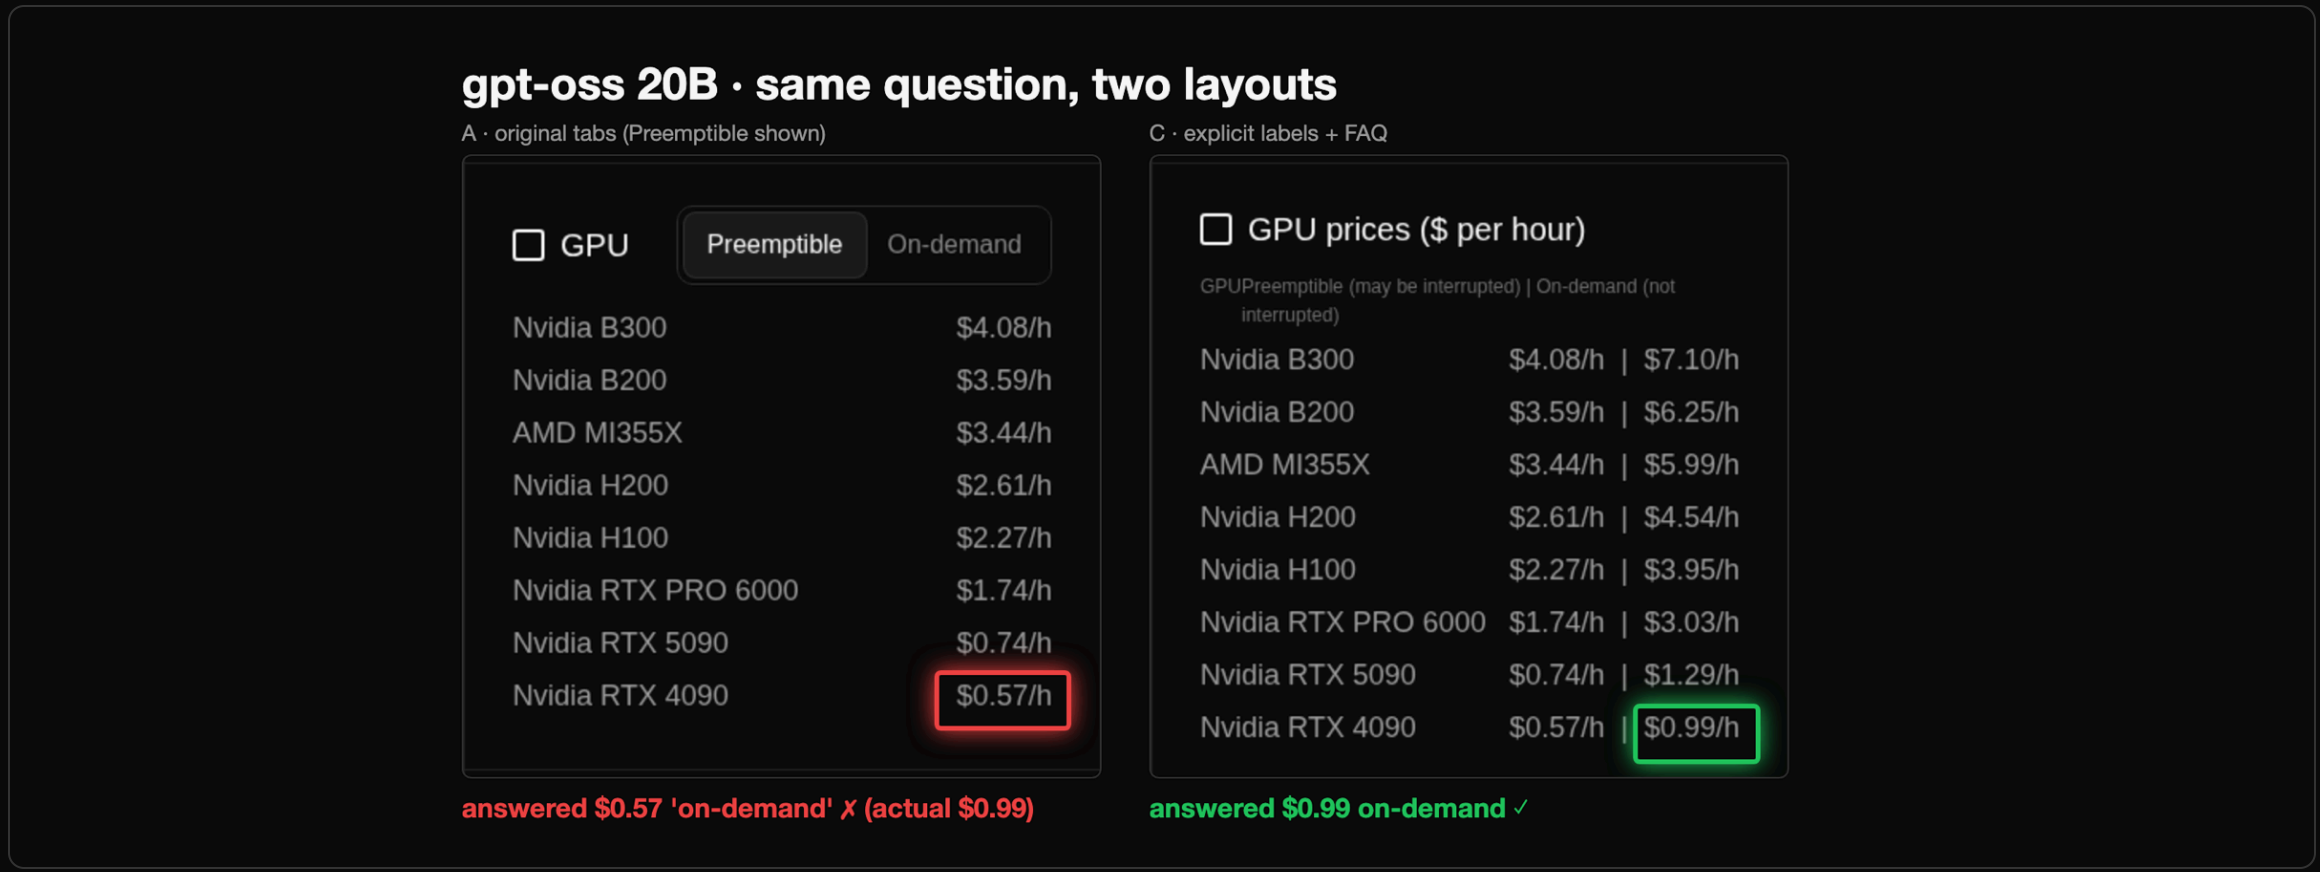

Human: the pill toggle shows which price is which. Screenshot from Chromium in the Daytona sandbox.

```
GPU
Preemptible
On-demand
Nvidia B300
$4.08/h
Nvidia B200
$3.59/h
AMD MI355X
$3.44/h
Nvidia H200
$2.61/h
Nvidia H100
$2.27/h
Nvidia RTX PRO 6000
$1.74/h
Nvidia RTX 5090
$0.74/h
Nvidia RTX 4090
$0.57/h
```

AI (gpt-oss 20B): the page's inner_text. "Preemptible" and "On-demand" are just two words; nothing says which list follows. It answered **$0.57 "on_demand"**.

Same run: 163227-9cca-gpt-oss-20b-A_original-q1-r0 (results/runs_rendered_v2_n10.jsonl, frames in results/frames/). Layout A reproduces the real page structure; it is not the live site.

# 60 real-browser runs in Daytona

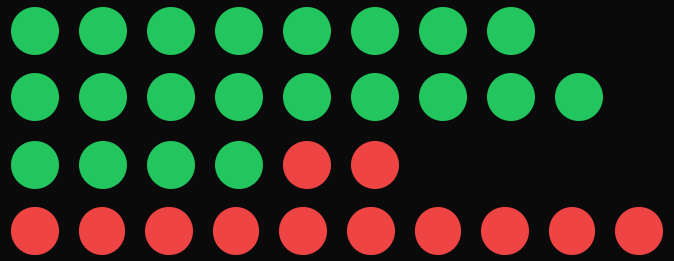



results/stage.html (40 s autoplay). Each run: a Daytona sandbox serves the layout, headless Chromium renders it, the reader can `click(text)`, every step is screenshotted.

60 runs in runs_rendered_v2_n10: A 10 reps × 4 models, C 5 reps × 4 models. 7 excluded as infra (browser/exec errors, LOG.md); 53 graded.

# Click the tab, or quote the wrong price

## Layout A · original tabs




correct 8/8 · clicked 8/8 · confident-wrong 0/8

correct 9/9 · clicked 9/9 · confident-wrong 0/9

correct 4/6 · clicked 6/6 · confident-wrong 2/6

correct 0/10 · clicked 0/10 · confident-wrong 10/10

Clicked → correct **21/23** · no click → correct **0/10**

**20/20**

layout C (both prices labelled + FAQ + JSON-LD), all 4 models

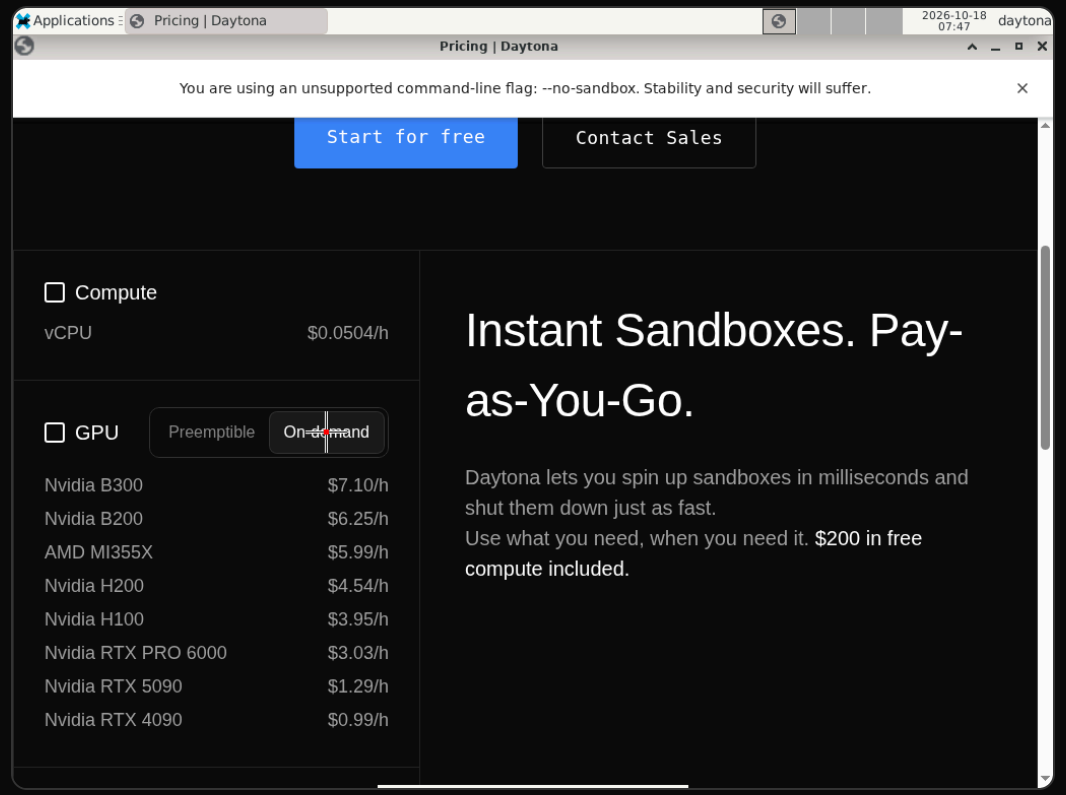
**0/12**

No page (answer from memory), q1 fully correct = 0%

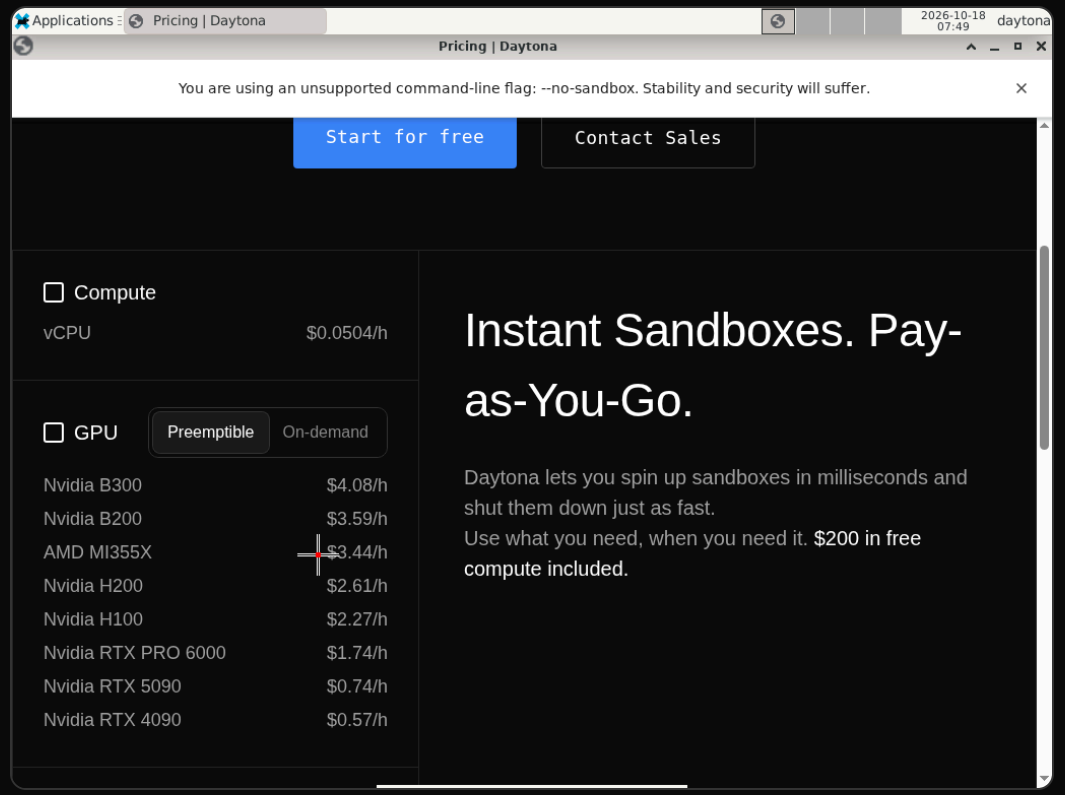
● correct ● confident wrong ● not found

Dots = graded runs, q1 (runs_rendered_v2_n10; No page from results/summary.json cells_by_question, mode none). Infra-excluded runs not shown.

# Computer use: pixels only qualitative, n=1




Claude · layout A · step 3: claude-sonnet-4-5: 1 click(s) at (308,417); answer 0.99 · verdict correct

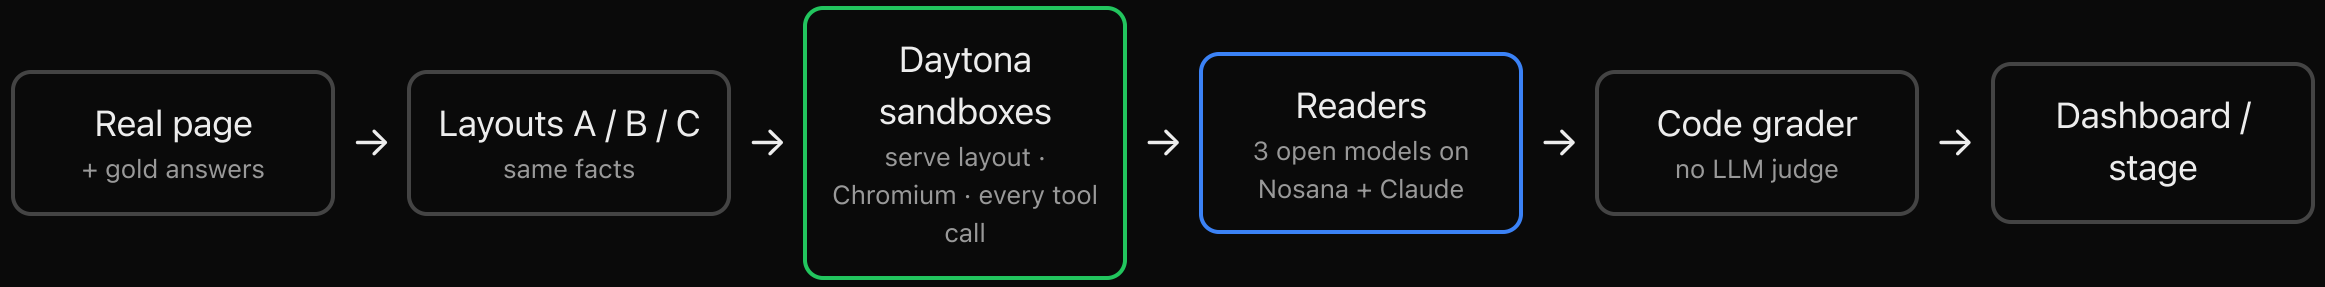



Gemma 4 26B · layout A · step 3: gemma4:26b: 6 click(s) at (301,538); answer none · verdict not_found

Daytona computer-use desktop + recording API. Screenshot in, `click(x,y)` out. Claude hit the On-demand pill; Gemma said "switch to On-demand" but clicked the MI355X row. Layout C: claude-sonnet-4-5: 0 click(s) at –; answer 0.99 · verdict correct. Coordinate miss may partly be Ollama image resizing (**unverified**).

results/cu_aside/ (run.json, step PNGs, recording.mp4). One run per cell, not part of the headline numbers.

# Pipeline




**32**

Daytona sandboxes

**741**

tool calls inside sandboxes (median 0.32s)

**391**

graded runs on Nosana GPUs (Qwen 3.5 9B, gpt-oss 20B, Gemma 4 26B)

Counts from results/runs*.jsonl, infra-excluded runs removed. 521 graded runs in total (results/summary.json).

# Product, next steps, limits

### Product

- Point it at your page: every AI reader × every layout, graded against your own answer key
- Shows which fact breaks and for which model, with replayable screenshots

### Next

- Commerce pages, e.g. Naver Smart Store product pages: options and prices behind tabs and selectors
- CI check: "does this layout change break any AI reader?"
- Auto-fix suggestion per failing fact

### Limits

- One page, reproduced layouts (not the live site), q1 for the browser runs
- Small n: A 10 reps/model, C 5 reps/model; directional, not proof
- Computer use is n=1
- OpenAI excluded (no API credits)

`github.com/TaeWoo1/daytona-nosana-hackathon · solo build, GitHub @TaeWoo1`

Appendix

# Judging criteria

## Completeness

- End-to-end: real page → gold → layouts → Daytona sandboxes → reader agents → code grader → dashboard, stage, deck
- 521 graded runs committed in results/, so every number can be checked without rerunning

## Innovation

- Measures comprehension accuracy (does the agent get the facts right?), not search visibility
- Reading mode as a dimension: crawler text, static visible text, real browser with click, computer use (n=1)

## Real-Life Problem Solving

- Layout reproducing daytona.io/pricing: runs with no click were correct 0/10; 12/33 runs confident wrong
- Same facts, explicit layout C: 20/20 correct

## Sponsored Product Usage

- Daytona: 32 ephemeral sandboxes; Chromium + every tool call inside the sandbox; computer-use desktop + recording
- Nosana: 3 open-weight models (Qwen 3.5 9B, gpt-oss 20B, Gemma 4 26B) on RTX 3090/4090

Rendered-browser numbers: results/runs_rendered_v2_n10.jsonl, q1, from summary.json → behavior.

Appendix

# q1 · visible_text

| | Qwen 3.5 9B | gpt-oss 20B | Gemma 4 26B | Claude | Worst model |
|---|---|---|---|---|---|
| A · original (tabs) | 0/3<br>wrong 1 | 0/3<br>wrong 3 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0%<br>Qwen 3.5 9B |
| B · explicit table | 1/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 33%<br>Qwen 3.5 9B |
| C · + FAQ + JSON-LD | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 100%<br>Qwen 3.5 9B |
| No page (memory) | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0%<br>Qwen 3.5 9B |

Cell = fully correct runs / runs (gold $0.99/h on-demand); "wrong" = runs with a confident wrong value. Source: results/summary.json (cells_by_question).

Appendix

# q1 · rendered_text_v2

| | Qwen 3.5 9B | gpt-oss 20B | Gemma 4 26B | Claude | Worst model |
|---|---|---|---|---|---|
| A · original (tabs) | 4/6<br>wrong 2 | 0/10<br>wrong 10 | 9/9<br>wrong 0 | 8/8<br>wrong 0 | 0%<br>gpt-oss 20B |
| C · + FAQ + JSON-LD | 5/5<br>wrong 0 | 5/5<br>wrong 0 | 5/5<br>wrong 0 | 5/5<br>wrong 0 | 100%<br>Qwen 3.5 9B |
| No page (memory) | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0%<br>Qwen 3.5 9B |

Cell = fully correct runs / runs (gold $0.99/h on-demand); "wrong" = runs with a confident wrong value. Source: results/summary.json (cells_by_question).

Appendix

# q1 · html_text

| | Qwen 3.5 9B | gpt-oss 20B | Gemma 4 26B | Claude | Worst model |
|---|---|---|---|---|---|
| A · original (tabs) | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 100%<br>Qwen 3.5 9B |
| B · explicit table | 0/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 0%<br>Qwen 3.5 9B |
| C · + FAQ + JSON-LD | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 3/3<br>wrong 0 | 100%<br>Qwen 3.5 9B |
| No page (memory) | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0/3<br>wrong 0 | 0%<br>Qwen 3.5 9B |

Cell = fully correct runs / runs (gold $0.99/h on-demand); "wrong" = runs with a confident wrong value. Source: results/summary.json (cells_by_question).

Appendix

# All 5 questions · html_text

| | Qwen 3.5 9B | gpt-oss 20B | Gemma 4 26B | Claude | Worst model |
|---|---|---|---|---|---|
| A · original (tabs) | 73% n=15 | 100% n=15 | 80% n=15 | 100% n=15 | 73% Qwen 3.5 9B |
| B · explicit table | 47% n=15 | 100% n=15 | 80% n=15 | 100% n=15 | 47% Qwen 3.5 9B |
| C · + FAQ + JSON-LD | 80% n=15 | 100% n=15 | 100% n=15 | 100% n=15 | 80% Qwen 3.5 9B |
| No page (memory) | 0% n=15 | 3% n=15 | 0% n=15 | 20% n=15 | 0% Gemma 4 26B |

Cell = mean share of answer fields matching gold. n = runs (5 questions × 3 reps). Outlined = weakest model per layout. Source: results/summary.json.

Appendix

# All 5 questions · visible_text

| | Qwen 3.5 9B | gpt-oss 20B | Gemma 4 26B | Claude | Worst model |
|---|---|---|---|---|---|
| A · original (tabs) | 80% n=15 | 90% n=15 | 67% n=15 | 90% n=15 | 67% Gemma 4 26B |
| B · explicit table | 60% n=15 | 100% n=15 | 80% n=15 | 93% n=15 | 60% Qwen 3.5 9B |
| C · + FAQ + JSON-LD | 80% n=15 | 100% n=15 | 100% n=15 | 100% n=15 | 80% Qwen 3.5 9B |
| No page (memory) | 0% n=15 | 3% n=15 | 0% n=15 | 20% n=15 | 0% Gemma 4 26B |

Cell = mean share of answer fields matching gold. n = runs (5 questions × 3 reps). Outlined = weakest model per layout. Source: results/summary.json.

Appendix

# Limitations

- One page; layouts are our reproductions of its structure (v1 plain mirror, v2 styled like the real page), not the live site.
- visible_text is a static approximation. Real-browser runs: q1 only, layouts A and C.
- Field accuracy gives partial credit; see the "fully correct" tables.
- Excluded: {"openai: RateLimitError": 60, "claude: BrowserRenderError": 2, "qwen3.5-9b: BrowserRenderError": 3, "qwen3.5-9b: DaytonaProcessExecutionTimeoutError": 1, "gemma4-26b: DaytonaProcessExecutionTimeoutError": 1}.
- Gold ambiguity: "first 5 GiB free" per sandbox or per org? (Both inside the q2 tolerance.)